

```

# IMPORTANT: RUN THIS CELL IN ORDER TO IMPORT YOUR KAGGLE DATA SOURCES
# TO THE CORRECT LOCATION (/kaggle/input) IN YOUR NOTEBOOK,
# THEN FEEL FREE TO DELETE THIS CELL.
# NOTE: THIS NOTEBOOK ENVIRONMENT DIFFERS FROM KAGGLE'S PYTHON
# ENVIRONMENT SO THERE MAY BE MISSING LIBRARIES USED BY YOUR
# NOTEBOOK.

import os
import sys
from tempfile import NamedTemporaryFile
from urllib.request import urlopen
from urllib.parse import unquote, urlparse
from urllib.error import HTTPError
from zipfile import ZipFile
import tarfile
import shutil

CHUNK_SIZE = 40960
DATA_SOURCE_MAPPING = ':https%3A%2F%2Fstorage.googleapis.com%2Fkaggle-data-sets%2F310%2F23498%2Fbundle%2Farchive.zip%3FX-Goog-Algorithm%3DZstd%3Fauth=1'

KAGGLE_INPUT_PATH='/kaggle/input'
KAGGLE_WORKING_PATH='/kaggle/working'
KAGGLE_SYMLINK='kaggle'

!umount /kaggle/input/ 2> /dev/null
shutil.rmtree('/kaggle/input', ignore_errors=True)
os.makedirs(KAGGLE_INPUT_PATH, 0o777, exist_ok=True)
os.makedirs(KAGGLE_WORKING_PATH, 0o777, exist_ok=True)

try:
    os.symlink(KAGGLE_INPUT_PATH, os.path.join(".", 'input'), target_is_directory=True)
except FileExistsError:
    pass
try:
    os.symlink(KAGGLE_WORKING_PATH, os.path.join(".", 'working'), target_is_directory=True)
except FileExistsError:
    pass

for data_source_mapping in DATA_SOURCE_MAPPING.split(','):
    directory, download_url_encoded = data_source_mapping.split(':')
    download_url = unquote(download_url_encoded)
    filename = urlparse(download_url).path
    destination_path = os.path.join(KAGGLE_INPUT_PATH, directory)
    try:
        with urlopen(download_url) as fileres, NamedTemporaryFile() as tfile:
            total_length = fileres.headers['content-length']
            print(f'Downloading {directory}, {total_length} bytes compressed')
            dl = 0
            data = fileres.read(CHUNK_SIZE)
            while len(data) > 0:
                dl += len(data)
                tfile.write(data)
                done = int(50 * dl / int(total_length))
                sys.stdout.write(f"\r[{ '=' * done }{' ' * (50-done)}] {dl} bytes downloaded")
                sys.stdout.flush()
                data = fileres.read(CHUNK_SIZE)
            if filename.endswith('.zip'):
                with ZipFile(tfile) as zfile:
                    zfile.extractall(destination_path)
            else:
                with tarfile.open(tfile.name) as tarfile:
                    tarfile.extractall(destination_path)
            print(f'\nDownloaded and uncompressed: {directory}')
    except HTTPError as e:
        print(f'Failed to load (likely expired) {download_url} to path {destination_path}')
        continue
    except OSError as e:
        print(f'Failed to load {download_url} to path {destination_path}')
        continue

print('Data source import complete.')

Downloading , 69155672 bytes compressed
[=====] 69155672 bytes downloaded
Downloaded and uncompressed:
Data source import complete.

```

```

# read & manipulate data
import pandas as pd
import numpy as np
import tensorflow as tf

# visualisations
import matplotlib.pyplot as plt
import seaborn as sns
sns.set(style='whitegrid', context='notebook')
%matplotlib notebook

# misc
import random as rn

# load the dataset
df = pd.read_csv('../input/creditcard.csv')

# manual parameters
RANDOM_SEED = 42
TRAINING_SAMPLE = 200000
VALIDATE_SIZE = 0.2

# setting random seeds for libraries to ensure reproducibility
np.random.seed(RANDOM_SEED)
rn.seed(RANDOM_SEED)
tf.random.set_seed(RANDOM_SEED) # Use tf.random.set_seed() instead of tf.set_random_seed()

# let's quickly convert the columns to lower case and rename the Class column
# so as to not cause syntax errors
df.columns = map(str.lower, df.columns)
df.rename(columns={'class': 'label'}, inplace=True)

# print first 5 rows to get an initial impression of the data we're dealing with
df.head()

```

	time	v1	v2	v3	v4	v5	v6	v7	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.0986
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.0851
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.2476
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.3774
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.2705

5 rows × 31 columns

```

# add a negligible amount to avoid taking the log of 0
df['log10_amount'] = np.log10(df.amount + 0.00001)

# keep the label field at the back
df = df[
    [col for col in df if col not in ['label', 'log10_amount']] +
    ['log10_amount', 'label']
]

```

```

# manual parameter
RATIO_TO_FRAUD = 15

# dropping redundant columns
df = df.drop(['time', 'amount'], axis=1)

# splitting by class
fraud = df[df.label == 1]
clean = df[df.label == 0]

# undersample clean transactions
clean_undersampled = clean.sample(
    int(len(fraud) * RATIO_TO_FRAUD),
    random_state=RANDOM_SEED
)

# concatenate with fraud transactions into a single dataframe
visualisation_initial = pd.concat([fraud, clean_undersampled])
column_names = list(visualisation_initial.drop('label', axis=1).columns)

# isolate features from labels
features, labels = visualisation_initial.drop('label', axis=1).values, \
    visualisation_initial.label.values

print(f""The non-fraud dataset has been undersampled from {len(clean):,} to {len(clean_undersampled):,}.
This represents a ratio of {RATIO_TO_FRAUD}:1 to fraud.""")

    The non-fraud dataset has been undersampled from 284,315 to 7,380.
    This represents a ratio of 15:1 to fraud.

from sklearn.manifold import TSNE
from mpl_toolkits.mplot3d import Axes3D

def tsne_scatter(features, labels, dimensions=2, save_as='graph.png'):
    if dimensions not in (2, 3):
        raise ValueError('tsne_scatter can only plot in 2d or 3d (what are you? An alien that can visualise >3d?). Make sure the "dimensi

    # t-SNE dimensionality reduction
    features_embedded = TSNE(n_components=dimensions, random_state=RANDOM_SEED).fit_transform(features)

    # initialising the plot
    fig, ax = plt.subplots(figsize=(8,8))

    # counting dimensions
    if dimensions == 3: ax = fig.add_subplot(111, projection='3d')

    # plotting data
    ax.scatter(
        *zip(*features_embedded[np.where(labels==1)]),
        marker='o',
        color='r',
        s=2,
        alpha=0.7,
        label='Fraud'
    )
    ax.scatter(
        *zip(*features_embedded[np.where(labels==0)]),
        marker='o',
        color='g',
        s=2,
        alpha=0.3,
        label='Clean'
    )

    # storing it to be displayed later
    plt.legend(loc='best')
    plt.savefig(save_as);
    plt.show;

tsne_scatter(features, labels, dimensions=2, save_as='tsne_initial_2d.png')

print(f""Shape of the datasets:
    clean (rows, cols) = {clean.shape}
    fraud (rows, cols) = {fraud.shape}""")

    Shape of the datasets:
    clean (rows, cols) = (284315, 30)
    fraud (rows, cols) = (492, 30)

```

```

# shuffle our training set
clean = clean.sample(frac=1).reset_index(drop=True)

# training set: exclusively non-fraud transactions
X_train = clean.iloc[:TRAINING_SAMPLE].drop('label', axis=1)

# testing set: the remaining non-fraud + all the fraud
X_test = clean.iloc[TRAINING_SAMPLE:].append(fraud).sample(frac=1)

<ipython-input-17-334e732d32cd>:8: FutureWarning: The frame.append method is deprecate
X_test = clean.iloc[TRAINING_SAMPLE:].append(fraud).sample(frac=1)

```

```

print(f""Our testing set is composed as follows:

{X_test.label.value_counts()}""")

Our testing set is composed as follows:

0      84315
1         492
Name: label, dtype: int64

from sklearn.model_selection import train_test_split

# train // validate - no labels since they're all clean anyway
X_train, X_validate = train_test_split(X_train,
                                       test_size=VALIDATE_SIZE,
                                       random_state=RANDOM_SEED)

# manually splitting the labels from the test df
X_test, y_test = X_test.drop('label', axis=1).values, X_test.label.values

print(f""Shape of the datasets:
training (rows, cols) = {X_train.shape}
validate (rows, cols) = {X_validate.shape}
holdout (rows, cols) = {X_test.shape}""")

Shape of the datasets:
training (rows, cols) = (160000, 29)
validate (rows, cols) = (40000, 29)
holdout (rows, cols) = (84807, 29)

from sklearn.preprocessing import Normalizer, MinMaxScaler
from sklearn.pipeline import Pipeline

# configure our pipeline
pipeline = Pipeline([('normalizer', Normalizer()),
                     ('scaler', MinMaxScaler())])

# get normalization parameters by fitting to the training data
pipeline.fit(X_train);

# transform the training and validation data with these parameters
X_train_transformed = pipeline.transform(X_train)
X_validate_transformed = pipeline.transform(X_validate)

g = sns.PairGrid(X_train.iloc[:, :3].sample(600, random_state=RANDOM_SEED))
plt.subplots_adjust(top=0.9)
g.fig.suptitle('Before:')
g.map_diag(sns.kdeplot)
g.map_offdiag(sns.kdeplot);

g = sns.PairGrid(pd.DataFrame(X_train_transformed, columns=column_names).iloc[:, :3].sample(600, random_state=RANDOM_SEED))
plt.subplots_adjust(top=0.9)
g.fig.suptitle('After:')
g.map_diag(sns.kdeplot)
g.map_offdiag(sns.kdeplot);

# Load the extension and start TensorBoard
%load_ext tensorboard
%tensorboard --logdir logs

```

TensorBoard

INACTIVE
SCALARS

☐ Show data download links
☐ Ignore outliers in chart scaling

Tooltip sorting method:

default ▼

Smoothing

0.6

Horizontal Axis

STEP RELATIVE

WALL

Runs

Write a regex to filter runs

TOGGLE ALL RUNS

logs

No scalar data was found.

Probable causes:

- You haven't written any scalar data to your event files.
- TensorBoard can't find your event files.

If you're new to using TensorBoard, and want to find out how to add data and set up your event files, check out the [README](#) and perhaps the [TensorBoard tutorial](#).

If you think TensorBoard is configured properly, please see [the section of the README devoted to missing data problems](#) and consider filing an issue on GitHub.

```
# data dimensions // hyperparameters
input_dim = X_train_transformed.shape[1]
BATCH_SIZE = 256
EPOCHS = 100

# https://keras.io/layers/core/
autoencoder = tf.keras.models.Sequential([

    # deconstruct / encode
    tf.keras.layers.Dense(input_dim, activation='elu', input_shape=(input_dim, )),
    tf.keras.layers.Dense(16, activation='elu'),
    tf.keras.layers.Dense(8, activation='elu'),
    tf.keras.layers.Dense(4, activation='elu'),
    tf.keras.layers.Dense(2, activation='elu'),

    # reconstruction / decode
    tf.keras.layers.Dense(4, activation='elu'),
    tf.keras.layers.Dense(8, activation='elu'),
    tf.keras.layers.Dense(16, activation='elu'),
    tf.keras.layers.Dense(input_dim, activation='elu')

])

# https://keras.io/api/models/model_training_apis/
autoencoder.compile(optimizer="adam",
                    loss="mse",
                    metrics=["acc"])

# print an overview of our model
autoencoder.summary();
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 29)	870
dense_1 (Dense)	(None, 16)	480
dense_2 (Dense)	(None, 8)	136
dense_3 (Dense)	(None, 4)	36
dense_4 (Dense)	(None, 2)	10
dense_5 (Dense)	(None, 4)	12
dense_6 (Dense)	(None, 8)	40
dense_7 (Dense)	(None, 16)	144
dense_8 (Dense)	(None, 29)	493
Total params: 2221 (8.68 KB)		
Trainable params: 2221 (8.68 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
from datetime import datetime

# current date and time
yyyymmddHHMM = datetime.now().strftime('%Y%m%d%H%M')

# new folder for a new run
log_subdir = f'{yyyymmddHHMM}_batch{BATCH_SIZE}_layers{len(autoencoder.layers)}'

# define our early stopping
early_stop = tf.keras.callbacks.EarlyStopping(
    monitor='val_loss',
    min_delta=0.0001,
    patience=10,
    verbose=1,
    mode='min',
    restore_best_weights=True
)

save_model = tf.keras.callbacks.ModelCheckpoint(
    filepath='autoencoder_best_weights.hdf5',
    save_best_only=True,
    monitor='val_loss',
    verbose=0,
    mode='min'
)

tensorboard = tf.keras.callbacks.TensorBoard(
    f'logs/{log_subdir}',
    batch_size=BATCH_SIZE,
    update_freq='batch'
)

# callbacks argument only takes a list
cb = [early_stop, save_model, tensorboard]
```



▼ Training

```
history = autoencoder.fit(
    X_train_transformed, X_train_transformed,
    shuffle=True,
    epochs=EPOCHS,
    batch_size=BATCH_SIZE,
    callbacks=cb,
    validation_data=(X_validate_transformed, X_validate_transformed)
);
```

Epoch 1/100
625/625 [=====] - 4s 6ms/step - loss: 0.0111 - acc: 0.373
Epoch 2/100
625/625 [=====] - 3s 5ms/step - loss: 0.0110 - acc: 0.377
Epoch 3/100
625/625 [=====] - 10s 16ms/step - loss: 0.0110 - acc: 0.377
Epoch 4/100
625/625 [=====] - 5s 8ms/step - loss: 0.0109 - acc: 0.380
Epoch 5/100
625/625 [=====] - 3s 5ms/step - loss: 0.0109 - acc: 0.381
Epoch 6/100
625/625 [=====] - 3s 5ms/step - loss: 0.0108 - acc: 0.381
Epoch 7/100
625/625 [=====] - 4s 6ms/step - loss: 0.0108 - acc: 0.381
Epoch 8/100
625/625 [=====] - 3s 5ms/step - loss: 0.0108 - acc: 0.381
Epoch 9/100
625/625 [=====] - 4s 7ms/step - loss: 0.0107 - acc: 0.382
Epoch 10/100
625/625 [=====] - 4s 6ms/step - loss: 0.0107 - acc: 0.384
Epoch 11/100
625/625 [=====] - 3s 5ms/step - loss: 0.0107 - acc: 0.385
Epoch 12/100
625/625 [=====] - 4s 6ms/step - loss: 0.0106 - acc: 0.386
Epoch 13/100
625/625 [=====] - 6s 10ms/step - loss: 0.0106 - acc: 0.386
Epoch 14/100
625/625 [=====] - 4s 7ms/step - loss: 0.0106 - acc: 0.387
Epoch 15/100
625/625 [=====] - 3s 5ms/step - loss: 0.0105 - acc: 0.387
Epoch 16/100
625/625 [=====] - 5s 9ms/step - loss: 0.0105 - acc: 0.387
Epoch 17/100
625/625 [=====] - 5s 8ms/step - loss: 0.0105 - acc: 0.388
Epoch 18/100
625/625 [=====] - 4s 6ms/step - loss: 0.0105 - acc: 0.388
Epoch 19/100
625/625 [=====] - 6s 10ms/step - loss: 0.0104 - acc: 0.388
Epoch 20/100
625/625 [=====] - 5s 9ms/step - loss: 0.0104 - acc: 0.389
Epoch 21/100
625/625 [=====] - 6s 10ms/step - loss: 0.0104 - acc: 0.390
Epoch 22/100
625/625 [=====] - 4s 6ms/step - loss: 0.0104 - acc: 0.391
Epoch 23/100
625/625 [=====] - 3s 5ms/step - loss: 0.0103 - acc: 0.392
Epoch 24/100
625/625 [=====] - 3s 5ms/step - loss: 0.0103 - acc: 0.393
Epoch 25/100
625/625 [=====] - 4s 6ms/step - loss: 0.0103 - acc: 0.394
Epoch 26/100
625/625 [=====] - 3s 5ms/step - loss: 0.0103 - acc: 0.396
Epoch 27/100
625/625 [=====] - 3s 5ms/step - loss: 0.0103 - acc: 0.397
Epoch 28/100
625/625 [=====] - 4s 6ms/step - loss: 0.0102 - acc: 0.399
Epoch 29/100
625/625 [=====] - 3s 5ms/step - loss: 0.0102 - acc: 0.400
Epoch 30/100
625/625 [=====] - 3s 5ms/step - loss: 0.0102 - acc: 0.401
Epoch 31/100
625/625 [=====] - 3s 5ms/step - loss: 0.0102 - acc: 0.402
Epoch 32/100
625/625 [=====] - 4s 6ms/step - loss: 0.0102 - acc: 0.403
Epoch 33/100
625/625 [=====] - 3s 5ms/step - loss: 0.0102 - acc: 0.404
Epoch 34/100
625/625 [=====] - 3s 5ms/step - loss: 0.0101 - acc: 0.404
Epoch 35/100
625/625 [=====] - 3s 5ms/step - loss: 0.0101 - acc: 0.405
Epoch 36/100
625/625 [=====] - 4s 7ms/step - loss: 0.0101 - acc: 0.406
Epoch 37/100
625/625 [=====] - 3s 5ms/step - loss: 0.0101 - acc: 0.408
Epoch 38/100
625/625 [=====] - 3s 5ms/step - loss: 0.0101 - acc: 0.407
Epoch 39/100
625/625 [=====] - 4s 6ms/step - loss: 0.0100 - acc: 0.409
Epoch 40/100
625/625 [=====] - 4s 6ms/step - loss: 0.0100 - acc: 0.410
Epoch 41/100
625/625 [=====] - 3s 5ms/step - loss: 0.0100 - acc: 0.411
Epoch 42/100
625/625 [=====] - 3s 5ms/step - loss: 0.0100 - acc: 0.413
Epoch 43/100
625/625 [=====] - 4s 6ms/step - loss: 0.0099 - acc: 0.413
Epoch 44/100
625/625 [=====] - 3s 5ms/step - loss: 0.0099 - acc: 0.415
Epoch 45/100
625/625 [=====] - 3s 5ms/step - loss: 0.0099 - acc: 0.415

```

Epoch 46/100
625/625 [=====] - 3s 5ms/step - loss: 0.0099 - acc: 0.416
Epoch 47/100
625/625 [=====] - 4s 7ms/step - loss: 0.0098 - acc: 0.417
Epoch 48/100
625/625 [=====] - 3s 5ms/step - loss: 0.0098 - acc: 0.418
Epoch 49/100
625/625 [=====] - 3s 5ms/step - loss: 0.0098 - acc: 0.417
Epoch 50/100
625/625 [=====] - 3s 5ms/step - loss: 0.0098 - acc: 0.419
Epoch 51/100
625/625 [=====] - 4s 6ms/step - loss: 0.0097 - acc: 0.419
Epoch 52/100
625/625 [=====] - 3s 5ms/step - loss: 0.0097 - acc: 0.420
Epoch 53/100
625/625 [=====] - 3s 5ms/step - loss: 0.0097 - acc: 0.422
Epoch 54/100
625/625 [=====] - 4s 7ms/step - loss: 0.0097 - acc: 0.423
Epoch 55/100
625/625 [=====] - 4s 6ms/step - loss: 0.0096 - acc: 0.424
Epoch 56/100
625/625 [=====] - 3s 5ms/step - loss: 0.0096 - acc: 0.423
Epoch 57/100
625/625 [=====] - 3s 5ms/step - loss: 0.0096 - acc: 0.424
Epoch 58/100
625/625 [=====] - 4s 7ms/step - loss: 0.0096 - acc: 0.425
Epoch 59/100
625/625 [=====] - 3s 5ms/step - loss: 0.0096 - acc: 0.425
Epoch 60/100
625/625 [=====] - 3s 5ms/step - loss: 0.0095 - acc: 0.425
Epoch 61/100
625/625 [=====] - 3s 5ms/step - loss: 0.0095 - acc: 0.425
Epoch 62/100
625/625 [=====] - 4s 6ms/step - loss: 0.0095 - acc: 0.425
Epoch 63/100
625/625 [=====] - 3s 5ms/step - loss: 0.0095 - acc: 0.425
Epoch 64/100
625/625 [=====] - 3s 5ms/step - loss: 0.0095 - acc: 0.428
Epoch 65/100
625/625 [=====] - 3s 5ms/step - loss: 0.0095 - acc: 0.427
Epoch 66/100
625/625 [=====] - 4s 6ms/step - loss: 0.0095 - acc: 0.428
Epoch 67/100
625/625 [=====] - 3s 5ms/step - loss: 0.0094 - acc: 0.428
Epoch 68/100
625/625 [=====] - 3s 5ms/step - loss: 0.0094 - acc: 0.429
Epoch 69/100
625/625 [=====] - 3s 5ms/step - loss: 0.0094 - acc: 0.429
Epoch 70/100
625/625 [=====] - 4s 6ms/step - loss: 0.0094 - acc: 0.429
Epoch 71/100
625/625 [=====] - 3s 5ms/step - loss: 0.0094 - acc: 0.429
Epoch 72/100
625/625 [=====] - 3s 5ms/step - loss: 0.0094 - acc: 0.430
Epoch 73/100
625/625 [=====] - 4s 6ms/step - loss: 0.0094 - acc: 0.430
Epoch 74/100
625/625 [=====] - 4s 6ms/step - loss: 0.0094 - acc: 0.430
Epoch 75/100
625/625 [=====] - 3s 5ms/step - loss: 0.0094 - acc: 0.430
Epoch 76/100
625/625 [=====] - 3s 5ms/step - loss: 0.0093 - acc: 0.431
Epoch 77/100
625/625 [=====] - 4s 7ms/step - loss: 0.0093 - acc: 0.430
Epoch 78/100
625/625 [=====] - 3s 5ms/step - loss: 0.0093 - acc: 0.431
Epoch 79/100
625/625 [=====] - 3s 5ms/step - loss: 0.0093 - acc: 0.431
Epoch 80/100
625/625 [=====] - 3s 5ms/step - loss: 0.0093 - acc: 0.431
Epoch 81/100
625/625 [=====] - 4s 7ms/step - loss: 0.0093 - acc: 0.431
Epoch 82/100
625/625 [=====] - 3s 5ms/step - loss: 0.0093 - acc: 0.431
Epoch 83/100
621/625 [=====>.] - ETA: 0s - loss: 0.0093 - acc: 0.4316Res
625/625 [=====] - 3s 5ms/step - loss: 0.0093 - acc: 0.431

```

```

# transform the test set with the pipeline fitted to the training set
X_test_transformed = pipeline.transform(X_test)

```

```

# pass the transformed test set through the autoencoder to get the reconstructed result
reconstructions = autoencoder.predict(X_test_transformed)

```



```

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not
warnings.warn(
Exception ignored in: <function _xla_gc_callback at 0x786b05726f80>
Traceback (most recent call last):
  File "/usr/local/lib/python3.10/dist-packages/jax/_src/lib/__init__.py", line 97, i
    def _xla_gc_callback(*args):
KeyboardInterrupt:
2651/2651 [=====] - 4s 1ms/step

```

```

# calculating the mean squared error reconstruction loss per row in the numpy array
mse = np.mean(np.power(X_test_transformed - reconstructions, 2), axis=1)

```

```

clean = mse[y_test==0]
fraud = mse[y_test==1]

```

```

fig, ax = plt.subplots(figsize=(6,6))

```

```

ax.hist(clean, bins=50, density=True, label="clean", alpha=.6, color="green")
ax.hist(fraud, bins=50, density=True, label="fraud", alpha=.6, color="red")

```

```

plt.title("(Normalized) Distribution of the Reconstruction Loss")
plt.legend()
plt.show()

```

```

THRESHOLD = 3

```

```

def mad_score(points):
    """https://www.itl.nist.gov/div898/handbook/eda/section3/eda35h.htm """
    m = np.median(points)
    ad = np.abs(points - m)
    mad = np.median(ad)

    return 0.6745 * ad / mad

```

```

z_scores = mad_score(mse)
outliers = z_scores > THRESHOLD

```

```

print(f"Detected {np.sum(outliers):,} outliers in a total of {np.size(z_scores):,} transactions [{np.sum(outliers)/np.size(z_scores):.2%}
      Detected 1,892 outliers in a total of 84,807 transactions [2.23%].

```

```

from sklearn.metrics import (confusion_matrix,
                             precision_recall_curve)

```

```

# get (mis)classification
cm = confusion_matrix(y_test, outliers)

```

```

# true/false positives/negatives
(tn, fp,
 fn, tp) = cm.flatten()

```

```

print(f""The classifications using the MAD method with threshold={THRESHOLD} are as follows:
{cm}

```

```

% of transactions labeled as fraud that were correct (precision): {tp}/{(fp)+(tp)} = {tp/(fp+tp):.2%}
% of fraudulent transactions were caught successfully (recall):    {tp}/{(fn)+(tp)} = {tp/(fn+tp):.2%}""")

```

```

The classifications using the MAD method with threshold=3 are as follows:
[[82784 1531]
 [ 131  361]]

```

```

% of transactions labeled as fraud that were correct (precision): 361/(1531+361) = 19

```

```

clean = z_scores[y_test==0]
fraud = z_scores[y_test==1]

```

```

fig, ax = plt.subplots(figsize=(6,6))

```

```

ax.hist(clean, bins=50, density=True, label="clean", alpha=.6, color="green")
ax.hist(fraud, bins=50, density=True, label="fraud", alpha=.6, color="red")

```

```

plt.title("Distribution of the modified z-scores")
plt.legend()
plt.show()

```

```
encoder = tf.keras.models.Sequential(autoencoder.layers[:5])
encoder.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 29)	870
dense_1 (Dense)	(None, 16)	480
dense_2 (Dense)	(None, 8)	136
dense_3 (Dense)	(None, 4)	36
dense_4 (Dense)	(None, 2)	10
Total params: 1532 (5.98 KB)		
Trainable params: 1532 (5.98 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
# taking all the fraud, undersampling clean
```

```
fraud = X_test_transformed[y_test==1]
```

```
clean = X_test_transformed[y_test==0][:len(fraud) * RATIO_TO_FRAUD, ]
```

```
# combining arrays & building labels
```

```
features = np.append(fraud, clean, axis=0)
```

```
labels = np.append(np.ones(len(fraud)),
                   np.zeros(len(clean)))
```

```
# getting latent space representation
```

```
latent_representation = encoder.predict(features)
```

```
print(f'Clean transactions downsampled from {len(X_test_transformed[y_test==0]):,} to {len(clean):,}.')
```

```
print('Shape of latent representation:', latent_representation.shape)
```

```
246/246 [=====] - 0s 1ms/step
```

```
Clean transactions downsampled from 84,315 to 7,380.
```